

no	Questions	Answers							
	1 What is Generative AI primarily designed to do?	Generate new content (text, image, code, etc.)							
	2 Which model architecture powers most modern LLMs?	Transformer							
	3 What does LLM stand for?	Large Language Model							
	4 What is "Prompt Engineering"?	Designing effective inputs for AI models							
	5 Which of the following is an example of a foundation model?	GPT Model							
	6 Fine-tuning a model means:	Adjusting pre-trained model on specific dataset							
	7 Token in LLM refers to:	Word or sub-word unit processed by model							
	8 Hallucination in Gen AI means:	Generating incorrect but confident answers							
	9 Which factor impacts LLM cost most?	Token usage							
	10 Responsible AI primarily focuses on:	Ethical and safe AI usage							
	11 Which AWS service provides foundation models?	Bedrock							
	12 RAG stands for:	Retrieval Augmented Generation							
	13 Embedding models are mainly used for:	Vector representation of text							
	14 Vector databases are used to:	Perform similarity search							
	15 Temperature parameter controls:	Creativity/randomness of output							
	16 Top-p sampling helps in:	Limiting randomness in output							
	17 Which is an example of a Gen AI use case?	Content generation							
	18 What is inference in AI?	Using trained model to generate output							
	19 Zero-shot prompting means:	No examples given to model							
	20 Few-shot prompting means:	Providing examples in prompt							
	21 Your company wants to build an AI chatbot using internal documents. What archi	Use RAG architecture with embeddings, a vector DB (like FAISS), and an LLM for generating responses based on relevant document retrieval.							
	22 How would you reduce hallucinations in an LLM-based system?	Reduce hallucinations by using RAG for fact-checking and fine-tuning the model on specific, high-quality datasets.							
	23 A client wants AI content generation but with strict compliance. What controls wo	Implement content filters, approval workflows, and audit trails to ensure compliance in AI-generated content.							
	24 How would you optimize cost in a Gen AI solution deployed in cloud?	Optimize costs by using serverless architecture, reserved instances, and scaling compute resources based on demand.							
	25 Explain the difference between fine-tuning and prompt engineering.	Fine-tuning adjusts a model's weights with specific data, while prompt engineering designs input text to guide the model's output.							
	26 Design a Gen AI solution for an automated customer email response system.	Design an email system that parses customer queries, retrieves info via vector DB, and generates responses with an LLM, using escalation when necessary.							
	27 If model responses are biased, how would you handle it?	Handle biased responses by curating diverse training data, using bias detection, and applying fairness filters.							
	28 How would you secure API access for LLM integration?	Secure API access with API keys, OAuth, rate limiting, and SSL/TLS encryption to protect the LLM integration.							